

Best AI Models

What to use in May 2026: coding, general reasoning, audio, and visual AI

closed + open-weight

best by task

snapshot date: 22 May 2026

General intelligence

Coding agents

Speech + music

Image + video

The May 2026 Landscape

No single winner: the best model depends on task, cost, deployment, and modality.

Closed frontier leads raw quality

Claude Opus 4.7 / Mythos, GPT-5.5, and Gemini 3.1 Pro split the top categories: coding depth, agentic terminal use, and multimodal reasoning.

Open-weight gap is shrinking

Kimi K2.6, DeepSeek V4, Qwen 3.6, GLM-5.x, Llama 4, and Mistral Large 3 are increasingly viable for self-hosting and cost control.

Multimodal becomes default

Top systems now mix text, code, image, audio, and video. The distinction is less “chatbot” vs “tool” and more “router over modalities.”

Decision rule

Use frontier closed models for hardest reasoning, coding, and reliability-sensitive work.
Use open-weight models when data control, self-hosting, latency, or unit economics matter.
Use specialized audio / image / video systems instead of one general model when media quality is the product.
Route across 2-4 models; avoid betting the whole stack on a single leaderboard winner.

Best General & Multi-Domain Models

The strongest general models are now differentiated by workflow, not just raw benchmark averages.

Claude Opus 4.7 / Mythos

Best for complex reasoning, long-context analysis, and hard software engineering. Mythos leads several provisional rankings but may be limited-access.

GPT-5.5

Best for agentic terminal workflows, computer use, broad ecosystem access, and productized assistants. Strong default where OpenAI tooling matters.

Gemini 3.1 Pro

Best value among US frontier models for multimodal, science/research, long context, and Google Cloud/RAG-heavy workflows.

Open-weight tier

Kimi K2.6, DeepSeek V4-Pro, Qwen 3.6, Llama 4, Mistral Large 3: pick by license, hosting budget, context, and language coverage.

Best practical stack: Claude for hard coding + Gemini for large multimodal context + GPT for terminal/product agents + one open-weight fallback.

Best Coding Models

Coding is split between real-world repo repair, terminal agents, and self-hostable code models.

Best raw coding

Claude Mythos Preview / Claude Opus 4.7
SWE-bench Verified and SWE-bench Pro
leadership; strong multi-file issue repair and
planning.

Best terminal agent

GPT-5.5 / Codex
Top choice when the task is shell-heavy, repo
navigation, debugging loops, and autonomous
tool use.

Best value frontier

Gemini 3.1 Pro
Strong coding plus very large
multimodal context at lower major-lab
pricing.

Best open-weight

Kimi K2.6 / DeepSeek V4-Pro / GLM-5.x
High coding scores with self-host or cheaper
hosted routes.

Coding picks by need

Enterprise copilot: Claude Opus 4.7 or GPT-5.5 depending on IDE/tooling fit.

Autonomous dev agent: GPT-5.5 for terminal-heavy work; Claude for patch quality and review depth.

Cost-sensitive code automation: Kimi K2.6 or DeepSeek V4-Pro behind a routing layer.

Do not rely only on HumanEval-style scores; use SWE-bench Pro, SWE-Rebench, Terminal-Bench, and repo-specific evals.

Open-Weight / “Free To Use” Shortlist

Free weights are not free operations: GPU, hosting, maintenance, and licenses still matter.

Kimi K2.6

Open-weight coding/value leader; strong agentic behavior

DeepSeek V4-Pro / Flash

Best cost-performance; Flash for very cheap throughput

Qwen 3.6 / Qwen-VL

Strong multilingual, coding, and open vision family

Llama 4 Scout / Maverick

Long-context open weights; good privacy/on-prem baseline

Mistral Large 3 / Small 4

EU-friendly open-weight options, Apache-style deployment appeal

Gemma 4

Smaller Apache-style models for commercial-safe local apps

Use open weights when you need data residency, predictable marginal cost, fine-tuning control, or offline/on-prem deployment.

Best Sound Models: Speech, TTS, Music

Audio is best treated as layers: STT, TTS, speech-to-speech, and music/song generation.

Speech-to-text

Deepgram Nova-3: production streaming default.
AssemblyAI Universal-2: transcription plus intelligence.
Whisper Large V3: open-source self-host baseline.
Google Chirp: broad-language batch accuracy.

Text-to-speech

Inworld Realtime TTS 1.5 Max and Gemini 3.1 Flash TTS lead arena-style rankings.
ElevenLabs v3: expressive quality and cloning.
Fish Audio S2 Pro / Kokoro: open-weight picks.

Native audio agents

GPT-Realtime-2 / GPT realtime family: speech-in, reasoning, speech-out without a separate STT/TTS chain.
Gemini Live and Grok/Step audio compete on speech reasoning.

Music + sound

Commercial: Suno, Udio, MiniMax/Mureka-style systems for polished songs.
Open: SongGeneration 2 / LeVo 2 and UniAudio 2 for music, sound, and audio tasks.

For voice products, choose latency first; for creator products, choose expressiveness and controllability first.

Best Visual Generation Models

Image and video generation are specialized markets: use leaderboards, but test your own prompt style.

Image generation

Top commercial families to evaluate: GPT-image / DALL-E line, Midjourney v7, Imagen 4, Ideogram, FLUX, and Stable Diffusion derivatives.
Best choice depends on realism, typography, editing, and brand consistency.

Video generation

Top commercial families to evaluate: Veo, Sora, Runway Gen-4, Kling, Luma Ray, Pika, and MiniMax video.
Key criteria: motion coherence, character consistency, audio, editing controls, and rights.

Open/local generation

FLUX / Stable Diffusion ecosystem for images; open video models are improving but still trail best closed systems for production polish.

Production rule: pick by output acceptance rate, not leaderboard rank alone.
Test with your real prompts: people, hands, text, product shots, diagrams, long camera moves, and brand constraints.
Track rights/licensing: training data, likeness, music/audio, and commercial usage terms vary widely.

Best Visual Understanding Models

Vision is now a core capability for documents, UI agents, video analysis, and multimodal RAG.

Best fine visual reasoning

Claude Opus 4.7
Strong document, chart, OCR-like, and fine-grained image analysis signals; good for review-heavy workflows.

Best broad multimodal

Gemini 3.1 Pro
Strong image + audio + video input with large context; natural fit for long video/document analysis.

Best UI/computer use

GPT-5.5
Strong ecosystem for computer-use, screen analysis, and agentic workflows where API/tooling integration matters.

Best open vision

Qwen 3.6-VL / Llama 4 vision-capable variants
Good self-host alternatives for OCR, VQA, screenshots, and private data.

Video understanding is a routing problem: combine keyframe extraction, transcript/audio, object/scene metadata, and a long-context multimodal model.

Recommended Model Routing

A practical “best models” answer is a portfolio, not a single model.

Hard coding + reviews

Claude Opus 4.7 / Mythos, with GPT-5.5 for terminal agents

Everyday assistant / agents

GPT-5.5 or Claude Sonnet/Opus depending on tool ecosystem

Research + multimodal docs

Gemini 3.1 Pro for long context and native audio/video/image input

Low-cost automation

DeepSeek V4 Flash / Pro, Kimi K2.6, Qwen family

Voice agents

Deepgram/Cartesia stack for latency; GPT-Realtime-2 when reasoning over speech matters

Image/video production

Specialized media models; evaluate per brand style and usage rights

Bottom line: keep a benchmark-informed shortlist, then run your own evals with real prompts, codebases, audio, images, and video.

Source Notes & Caveats

Benchmarks move quickly; this deck is a dated market snapshot.

Primary benchmark references used: SWE-bench / BenchLM / LLM Stats for coding; Artificial Analysis and public arenas for model/media comparisons.

News and synthesis references used: FutureAGI, BuildFastWithAI, o-mega, DeepLearning.AI, TTS Arena, and public GitHub/model-card pages.

Some leading systems are proprietary, API-only, region-limited, waitlisted, or only available to selected partners.

“Open source” is often actually “open weights”; verify license, commercial rights, acceptable-use policy, and deployment obligations.

Best practice: build a small internal eval suite before procurement or model migration.

Prepared for Mariusz | 22 May 2026